Clustering

Clustering

- Explore the structure of the dataset by splitting the data points into different groups via well-studied clustering algorithms

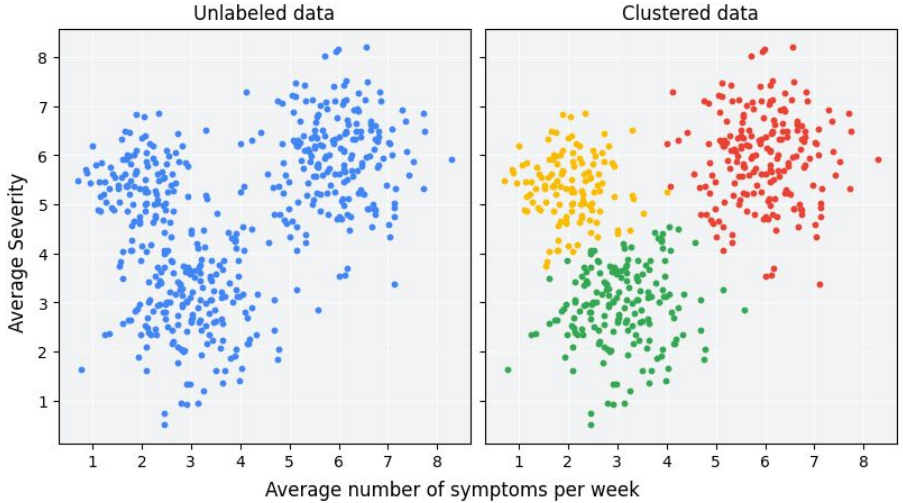

Image source: https://developers.google.com/machine-learning/clustering/overview

Clustering

- When some labels of a classification task are not available, clustering methods can sometimes be used to fill in the rest of the labels

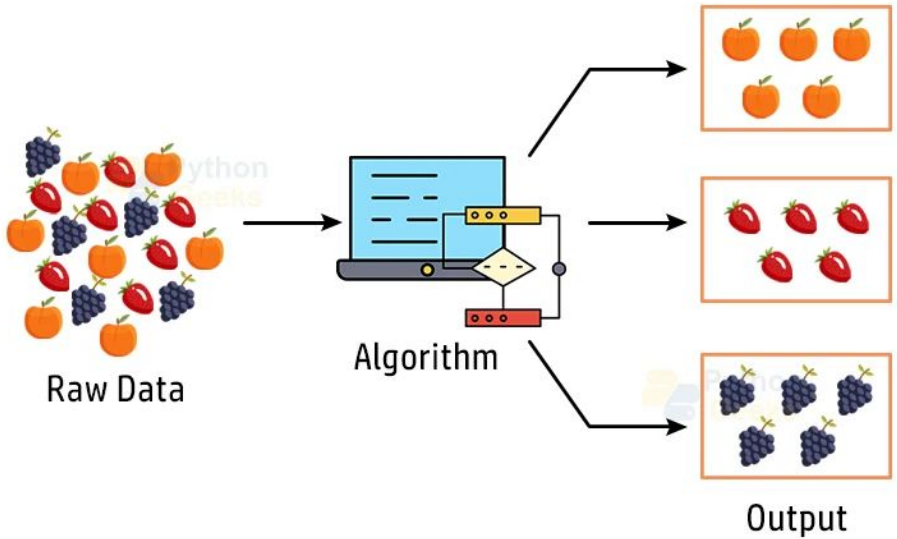

Image source: https://pythongeeks.org/clustering-in-machine-learning/

Clustering

- It can also be used to identify and visualize interesting groups of data
- If they are outliers in your dataset that form a cluster, you know there is probably a common source/reason for the outliers

Clustering

- Approximate idea: group geometrically close points into the same group, and geometrically distant points into different groups

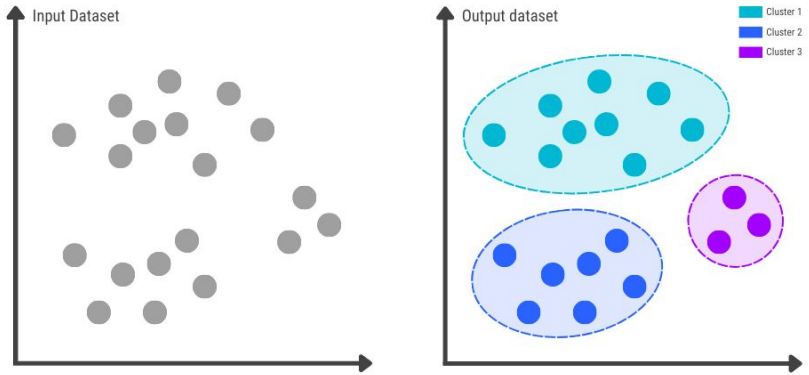

Image source: https://www.pi.exchange/blog/clustering-in-machine-learning

Clustering Algorithms

- Centroid-based: K-Means
- Density-based: DBSCAN family
- Gaussian mixture models (GMMs); they are useful for clustering data into visible clusters with complex shapes or even those that overlap

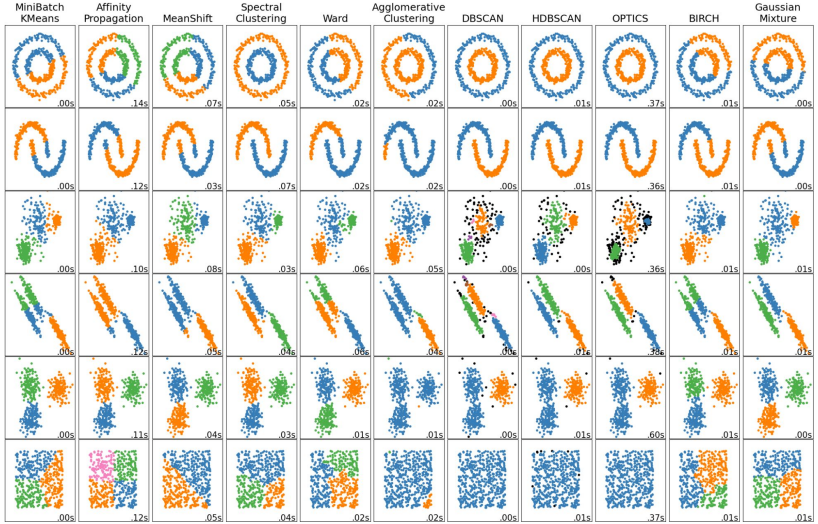

Image source: https://scikit-learn.org/stable/modules/clustering.html

K-Means Clustering

- Specify K random initial points as the initial “centroids” of the K clusters (that is where the “K” in the “K-Means” comes from)

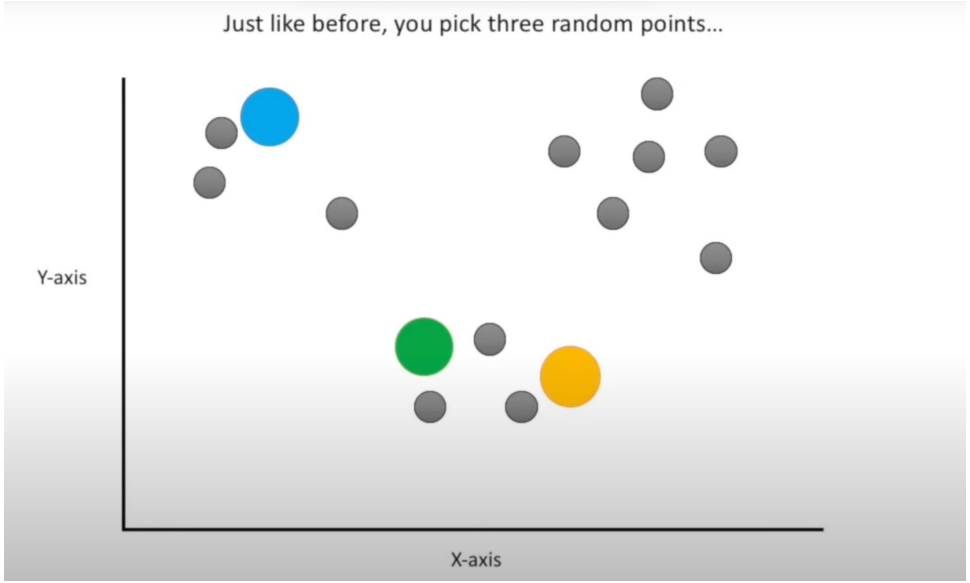

Image source: https://www.youtube.com/watch?v=4b5d3muPQmA

K-Means Clustering

Video: https://www.youtube.com/watch?v=4b5d3muPQmA

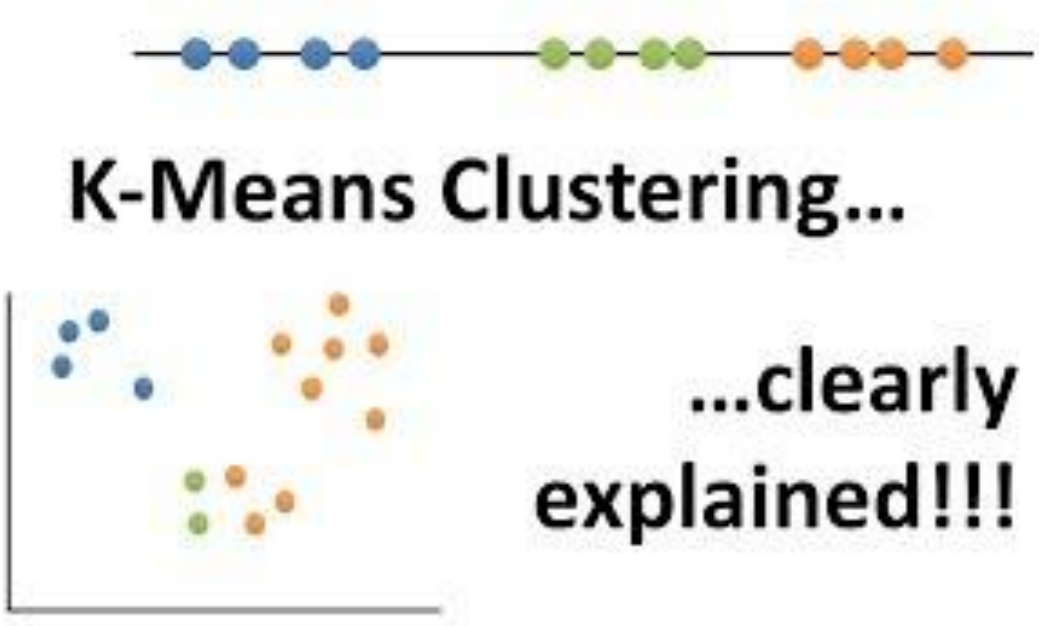

K-Means Clustering

- Associate each remaining data point with the centroid that is the least distant from that point (that is the initial clustering)

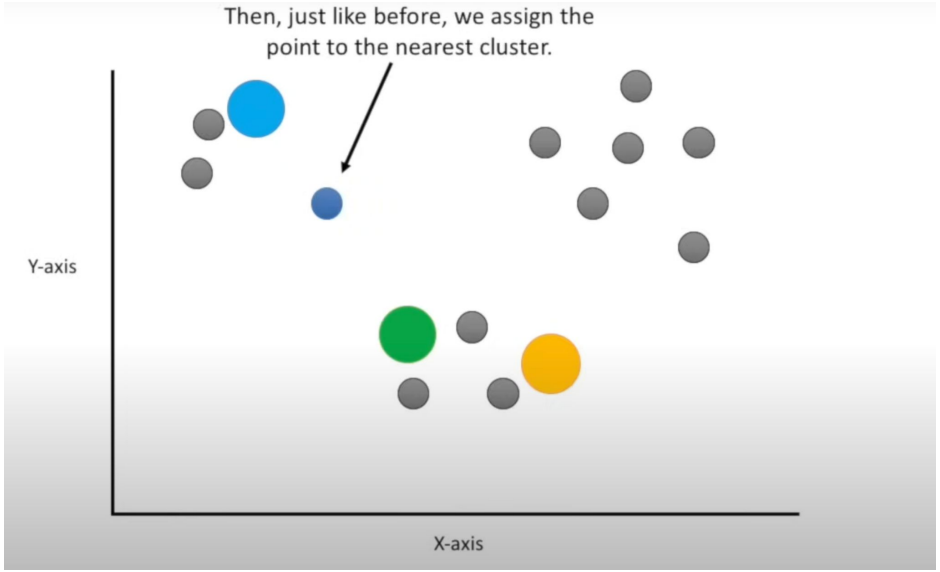

Image source: https://www.youtube.com/watch?v=4b5d3muPQmA

K-Means Clustering

- Recalculate the centroids based on the initial clustering, and then ignore all clustering labels from the data to cluster the data again based on the new centroids

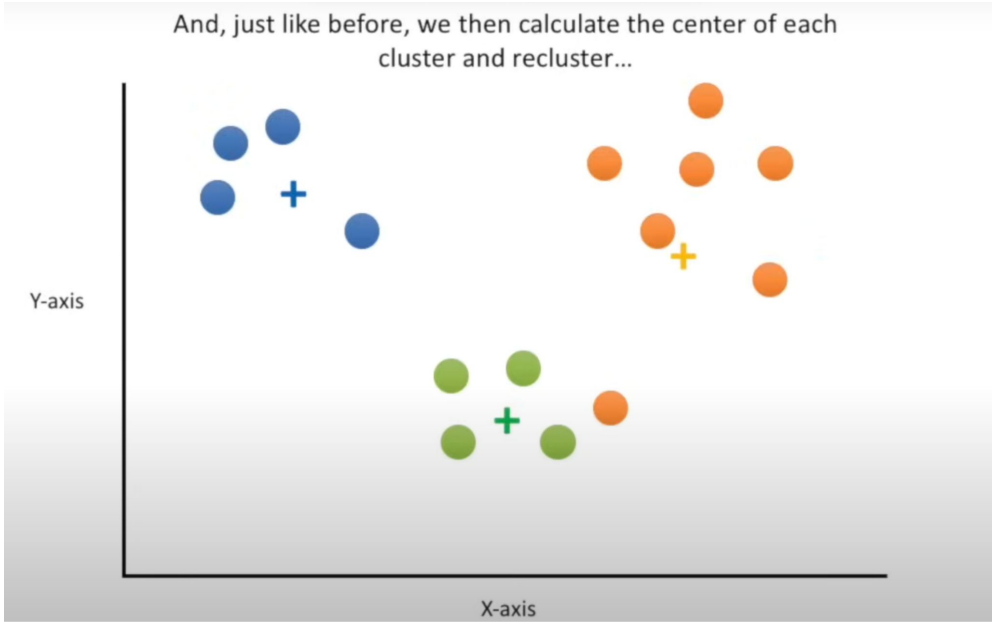

Image source: https://www.youtube.com/watch?v=4b5d3muPQmA

K-Means Clustering

- Iterative refinement: Repeat the process until it converges (i.e., nothing changes throughout cycles of reclustering) or it reaches the maximum step limit

K-Means Clustering

- Advantages: allows specifying the number of classes, which can be useful when you know what you need to classify data into
- Advantages: can classify unseen points as members of one of the clusters

K-Means Clustering

- Disadvantages: it cannot effectively model nonlinear dependencies, and may not work well for highly asymmetric data
- It relies on calculating mean values of points so it can be sensitive to outliers

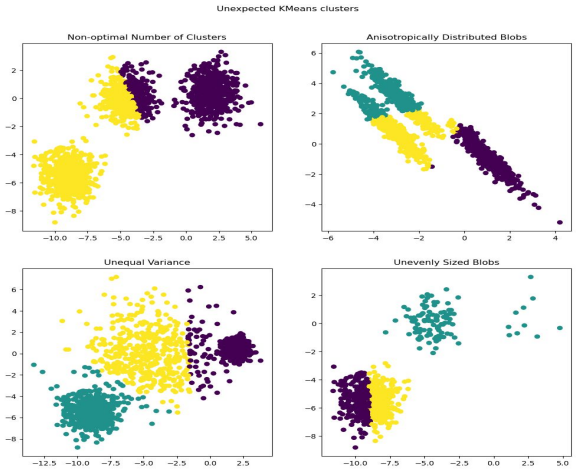

Image source: https://scikit-learn.org/stable/modules/clustering.html

DBSCAN

Video: https://www.youtube.com/watch?v=RDZUdRSDOok

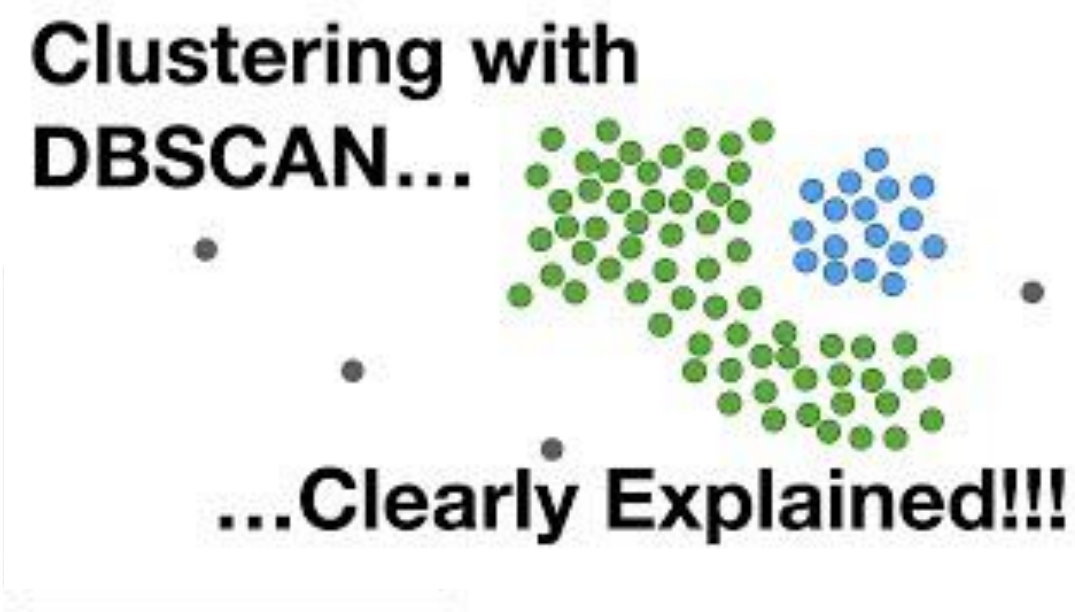

DBSCAN

- Specify: eps (how close must a point be to extend the cluster) and min_samples (how many points we need to make a cluster)

DBSCAN

- Randomly select a point; if a point has (min_samples - 1) neighbors within a distance of eps, it is selected as a “core point” and a cluster is formed
- Expand the cluster based on the distance threshold eps; if a point has enough close neighbors to start a cluster itself, it can add other points; otherwise it cannot add other points
- Repeat until no points can be added
- Find another core point from the remaining data, and do the cluster expansion process again. Repeat until no new clusters can be formed.
- Dense regions tend to get clustered together

(add youtube video screenshots every step of the way)

DBSCAN

- Advantages: can model nonlinear boundaries between classes
- Disadvantages: requires good understanding of the scale of the data and careful hyperparameter tuning

Try It Out

- Experience K-Means and DBSCAN clustering in this notebook: https://colab.research.google.com/drive/196UkbTkRnz2X7-CXSpHyCcFp3r9xHZio#scrollTo=6zf-9_wtTzVW

Useful Tip

- If you struggle to achieve the right separation, try to separate the dataset into more clusters
- And then group some of these clusters together to achieve the separation you want
- This has appeared in the Malaysia AI Olympiad (eye of feature engineering question) and is a decent-scoring solution in the IOAI (IOAI 2025 Q5) that doesn't use pseudo-labels to train a classifier.